

!!Capstone Project!!

BY

1. Vignya Durvasula 20P61A6615

2. Shreejit Cheela 20P61A6647

▼ Import Library

```
import pandas as pd#for data manipulation
import numpy as np#for working with arrays in the dataset
import seaborn as sns#for dtatistical data visualisation
import matplotlib.pyplot as plt#for dynamic data visualisation
```

▼ Import Dataset

```
df=pd.read_csv('https://github.com/ybifoundation/Dataset/blob/main/Melbourne%20Housing%20Market.csv?raw=true')#importing the dataset
```

```
df#displaying the dataframe
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	...	Bathroom	Car	Lands
0	Abbotsford	85 Turner St	2	h	1480000.0	S	Biggin	3/12/2016	2.5	3067.0	...	1.0	1.0	20
1	Abbotsford	25 Bloomburg St	2	h	1035000.0	S	Biggin	4/02/2016	2.5	3067.0	...	1.0	0.0	15
2	Abbotsford	5 Charles St	3	h	1465000.0	SP	Biggin	4/03/2017	2.5	3067.0	...	2.0	0.0	13
3	Abbotsford	40 Federation La	3	h	850000.0	PI	Biggin	4/03/2017	2.5	3067.0	...	2.0	1.0	9
4	Abbotsford	55a Park St	4	h	1600000.0	VB	Nelson	4/06/2016	2.5	3067.0	...	1.0	2.0	12
...
13575	Wheelers Hill	12 Strada Cr	4	h	1245000.0	S	Barry	26/08/2017	16.7	3150.0	...	2.0	2.0	65
13576	Williamstown	77 Merrett Dr	3	h	1031000.0	SP	Williams	26/08/2017	6.8	3016.0	...	2.0	2.0	33
13577	Williamstown	83 Power St	3	h	1170000.0	S	Raine	26/08/2017	6.8	3016.0	...	2.0	4.0	43
13578	Williamstown	96 Verdon St	4	h	2500000.0	PI	Sweeney	26/08/2017	6.8	3016.0	...	1.0	5.0	86

▼ Data Preprocessing

`df.shape` built-in function is used to know the total number of rows and columns in the dataset

```
(13580, 21)
```

```
df.nunique()#represents the number of unique values under each attribute
```

```
Suburb      314
Address     13378
Rooms        9
Type         3
Price       2204
Method        5
SellerG     268
Date         58
Distance    202
Postcode    198
Bedroom2     12
Bathroom     9
Car          11
Landsize    1448
BuildingArea 602
YearBuilt    144
CouncilArea   33
Lattitude    6503
Longitude    7063
Regionname    8
Propertycount 311
dtype: int64
```

```
df.columns#displaying column names
```

```
Index(['Suburb', 'Address', 'Rooms', 'Type', 'Price', 'Method', 'SellerG',
       'Date', 'Distance', 'Postcode', 'Bedroom2', 'Bathroom', 'Car',
       'Landsize', 'BuildingArea', 'YearBuilt', 'CouncilArea', 'Lattitude',
       'Longitude', 'Regionname', 'Propertycount'],
      dtype='object')
```

```
df.info()#displays the basic information of the dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13580 entries, 0 to 13579
```

```
Data columns (total 21 columns):
#      Column      Non-Null Count  Dtype
---  -
0      Suburb      13580 non-null  object
1      Address      13580 non-null  object
2      Rooms        13580 non-null  int64
3      Type         13580 non-null  object
4      Price        13580 non-null  float64
5      Method       13580 non-null  object
6      SellerG      13580 non-null  object
7      Date         13580 non-null  object
8      Distance     13580 non-null  float64
9      Postcode     13580 non-null  float64
10     Bedroom2     13580 non-null  float64
11     Bathroom     13580 non-null  float64
12     Car          13518 non-null  float64
13     Landsize     13580 non-null  float64
14     BuildingArea  7130 non-null   float64
15     YearBuilt     8205 non-null   float64
16     CouncilArea   12211 non-null  object
17     Lattitude     13580 non-null  float64
18     Longitude     13580 non-null  float64
19     Regionname    13580 non-null  object
20     Propertycount 13580 non-null  float64
dtypes: float64(12), int64(1), object(8)
memory usage: 2.2+ MB
```

```
df['Date']=pd.to_datetime(df['Date'])#converting the date attribute from object datatype to date datatype
```

▼ Adjusting Missing Values

```
df.drop('CouncilArea',axis='columns', inplace=True)#Because its of object datatype moreover we already many other attributes that ref
```

```
df['BuildingArea'].fillna(df.groupby(['Type'])['BuildingArea'].transform('mean'), inplace=True)#filling the missing values in a parti
df['YearBuilt'].fillna(df.groupby(['Type'])['YearBuilt'].transform('mean'), inplace=True)
```

```
df['Car'].fillna(df.groupby(['Type'])['Car'].transform('mean'), inplace=True)
```

```
df.info()#One can observe that the missing values are all now filled up
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 13580 entries, 0 to 13579
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Suburb           13580 non-null  object
1   Address          13580 non-null  object
2   Rooms            13580 non-null  int64
3   Type             13580 non-null  object
4   Price            13580 non-null  float64
5   Method           13580 non-null  object
6   SellerG          13580 non-null  object
7   Date             13580 non-null  datetime64[ns]
8   Distance         13580 non-null  float64
9   Postcode         13580 non-null  float64
10  Bedroom2         13580 non-null  float64
11  Bathroom         13580 non-null  float64
12  Car              13580 non-null  float64
13  Landsize         13580 non-null  float64
14  BuildingArea     13580 non-null  float64
15  YearBuilt        13580 non-null  float64
16  Lattitude        13580 non-null  float64
17  Longitude        13580 non-null  float64
18  Regionname       13580 non-null  object
19  Propertycount    13580 non-null  float64
dtypes: datetime64[ns](1), float64(12), int64(1), object(6)
memory usage: 2.1+ MB
```

```
df.describe()#displaying description of the data in the DataFrame
```

	Rooms	Price	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingAr
count	13580.000000	1.358000e+04	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000	13580.0000
mean	2.937997	1.075684e+06	10.137776	3105.301915	2.914728	1.534242	1.610716	558.416127	152.4892
std	0.955748	6.393107e+05	5.868725	90.676964	0.965921	0.691712	0.960511	3990.669241	392.9573
min	1.000000	8.500000e+04	0.000000	3000.000000	0.000000	0.000000	0.000000	0.000000	0.0000
25%	2.000000	6.500000e+05	6.100000	3044.000000	2.000000	1.000000	1.000000	177.000000	96.7500
50%	3.000000	9.030000e+05	9.200000	3084.000000	3.000000	1.000000	2.000000	440.000000	159.0000
75%	3.000000	1.330000e+06	13.000000	3148.000000	3.000000	2.000000	2.000000	651.000000	176.8662
max	10.000000	9.000000e+06	10.100000	3077.000000	30.000000	3.000000	10.000000	100014.000000	11515.0000

df.nunique()

Suburb	314
Address	13378
Rooms	9
Type	3
Price	2204
Method	5
SellerG	268
Date	58
Distance	202
Postcode	198
Bedroom2	12
Bathroom	9
Car	13
Landsize	1448
BuildingArea	605
YearBuilt	147
Lattitude	6503
Longtitude	7063
Regionname	8
Propertycount	311

dtype: int64

▼ Encoding

Converting the object or human readable values to numeric or machine readable form.

```
df['Method'].value_counts()#displaying the number of unique values under each label of the attribute
```

```
S      9022
SP     1703
PI     1564
VB     1199
SA       92
Name: Method, dtype: int64
```

```
df.replace({'Method':{'S':1,'SP':2,'PI':2,'VB':2,'SA':2}},inplace=True)#grouping
```

```
df['Method'].value_counts()#one can observe the change in label and value count
```

```
1      9022
2     4558
Name: Method, dtype: int64
```

```
df['Type'].value_counts()
```

```
h      9449
u      3017
t      1114
Name: Type, dtype: int64
```

```
df.replace({'Type':{'h':1,'u':2,'t':2}},inplace=True)
```

```
df['Type'].value_counts()
```

```
1    9449
2    4131
Name: Type, dtype: int64
```

```
df['Regionname'].value_counts()
```

```
Southern Metropolitan    4695
Northern Metropolitan    3890
Western Metropolitan     2948
Eastern Metropolitan     1471
South-Eastern Metropolitan  450
Eastern Victoria         53
Northern Victoria        41
Western Victoria         32
Name: Regionname, dtype: int64
```

```
df.replace({'Regionname':{'Southern Metropolitan':1,'Northern Metropolitan':2,'Western Metropolitan':3,'Eastern Metropolitan':3,'Sout
```

```
df['Regionname'].value_counts()
```

```
3    4995
1    4695
2    3890
Name: Regionname, dtype: int64
```

Explaining why these Attributes are dropped while defining our dependent and independent variable

```
df['SellerG'].value_counts()
```

```
Nelson    1565
Jellis    1316
hockingstuart 1167
```



```
Barry      1011
Ray        701
...
Prowse     1
Luxe       1
Zahn       1
Homes      1
Point      1
Name: SellerG, Length: 268, dtype: int64
```

```
df['Suburb'].value_counts()
```

```
Reservoir      359
Richmond       260
Bentleigh East  249
Preston        239
Brunswick      222
...
Sandhurst      1
Bullengarook   1
Croydon South  1
Montrose       1
Monbulk        1
Name: Suburb, Length: 314, dtype: int64
```

```
df['Address'].value_counts()
```

```
36 Aberfeldie St    3
2 Bruce St          3
5 Charles St         3
53 William St       3
14 Arthur St        3
..
16 Alleford St      1
2/1073 Centre Rd    1
14 Columbia St      1
21 Hardy Ct         1
6 Agnes St          1
Name: Address, Length: 13378, dtype: int64
```

Explanation:

Clearly, the lenght or the number of Categories of these Attributes is large and nearly impossible to Encode.

▼ Representating Encoded dataset and information

df

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	Bedroom2	Bathroom	Car	Lands
0	Abbotsford	85 Turner St	2	1	1480000.0	1	Biggin	2016-03-12	2.5	3067.0	2.0	1.0	1.0	20
1	Abbotsford	25 Bloomburg St	2	1	1035000.0	1	Biggin	2016-04-02	2.5	3067.0	2.0	1.0	0.0	10
2	Abbotsford	5 Charles St	3	1	1465000.0	2	Biggin	2017-04-03	2.5	3067.0	3.0	2.0	0.0	10
		40						2017-						

```
df.describe()
```

	Rooms	Type	Price	Method	Distance	Postcode	Bedroom2	Bathroom	Car
count	13580.000000	13580.000000	1.358000e+04	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000
mean	2.937997	1.304197	1.075684e+06	1.335641	10.137776	3105.301915	2.914728	1.534242	1.610716
std	0.955748	0.460084	6.393107e+05	0.472231	5.868725	90.676964	0.965921	0.691712	0.960511
min	1.000000	1.000000	8.500000e+04	1.000000	0.000000	3000.000000	0.000000	0.000000	0.000000
25%	2.000000	1.000000	6.500000e+05	1.000000	6.100000	3044.000000	2.000000	1.000000	1.000000
50%	3.000000	1.000000	9.030000e+05	1.000000	9.200000	3084.000000	3.000000	1.000000	2.000000
75%	3.000000	2.000000	1.330000e+06	2.000000	13.000000	3148.000000	3.000000	2.000000	2.000000
max	10.000000	2.000000	9.000000e+06	2.000000	48.100000	3977.000000	20.000000	8.000000	10.000000



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 13580 entries, 0 to 13579

Data columns (total 20 columns):

#	Column	Non-Null Count	Dtype
0	Suburb	13580 non-null	object
1	Address	13580 non-null	object
2	Rooms	13580 non-null	int64
3	Type	13580 non-null	int64
4	Price	13580 non-null	float64
5	Method	13580 non-null	int64
6	SellerG	13580 non-null	object
7	Date	13580 non-null	datetime64[ns]
8	Distance	13580 non-null	float64
9	Postcode	13580 non-null	float64
10	Bedroom2	13580 non-null	float64
11	Bathroom	13580 non-null	float64
12	Car	13580 non-null	float64
13	Landsize	13580 non-null	float64
14	BuildingArea	13580 non-null	float64
15	YearBuilt	13580 non-null	float64
16	Lattitude	13580 non-null	float64
17	Longitude	13580 non-null	float64
18	Regionname	13580 non-null	int64
19	Propertycount	13580 non-null	float64

dtypes: datetime64[ns](1), float64(12), int64(4), object(3)

memory usage: 2.1+ MB

df.nunique()

Suburb	314
Address	13378
Rooms	9
Type	2
Price	2204
Method	2
SellerG	268
Date	58
Distance	202
Postcode	198
Bedroom2	12

```

Bathroom      9
Car           13
Landsize      1448
BuildingArea   605
YearBuilt     147
Latitude      6503
Longitude     7063
Regionname     3
Propertycount 311
dtype: int64

```

```
df.groupby(['Type', 'Method']).count()#grouping categorical variables for feature selection
```

		Suburb	Address	Rooms	Price	SellerG	Date	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea
Type	Method													
1	1	6507	6507	6507	6507	6507	6507	6507	6507	6507	6507	6507	6507	6507
	2	2942	2942	2942	2942	2942	2942	2942	2942	2942	2942	2942	2942	2942
2	1	2515	2515	2515	2515	2515	2515	2515	2515	2515	2515	2515	2515	2515
	2	1616	1616	1616	1616	1616	1616	1616	1616	1616	1616	1616	1616	1616

▼ Data Visualisation

```
sns.pairplot(df)#visualisation study
```

Part 1: Regression Analysis

▼ Step 1: Define X and y

```
y=df['Price']# 'Price' is our Regression target  
np.log(y)#for increasing efficiency
```

```
0      14.207553  
1      13.849912  
2      14.197366  
3      13.652992  
4      14.285514  
...  
13575   14.034646  
13576   13.846040  
13577   13.972514  
13578   14.731801  
13579   14.066269  
Name: Price, Length: 13580, dtype: float64
```

```
y.shape
```

```
(13580,)
```

```
X=df.drop(['Suburb','Address','SellerG','Date', 'Price'],axis=1)#defining X
```

```
X.shape
```

```
(13580, 15)
```

▼ Dealing with Oversampling Data

```
'Price'
```

```
from imblearn.over_sampling import RandomOverSampler
```

```
r=RandomOverSampler(random_state=2408)
```

Before Oversampling

```
X.shape, y.shape
```

```
((13580, 15), (13580,))
```

```
y.value_counts()
```

```
1100000.0    113
1300000.0    109
650000.0     109
800000.0     109
600000.0     104
...
1928000.0      1
2236000.0      1
601500.0       1
550500.0       1
1323000.0      1
Name: Price, Length: 2204, dtype: int64
```

```
X.value_counts()
```

Rooms	Type	Method	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitude	R
2	2	1	3.3	3141.0	2.0	2.0	2.0	17200.0	80.737121	2000.000000	-37.83610	144.99660	1
1	2	1	3.3	3141.0	1.0	1.0	1.0	0.0	80.737121	2000.000000	-37.83610	144.99660	1
3	1	1	7.5	3123.0	3.0	2.0	2.0	431.0	120.000000	1950.000000	-37.82690	145.04960	1
		2	9.2	3146.0	3.0	2.0	1.0	704.0	134.000000	1940.000000	-37.85200	145.09420	1

2	2	1	6.5	3071.0	2.0	1.0	1.0	0.0	80.737121	1970.000000	-37.75610	145.00670	2
3	1	1	3.8	3207.0	3.0	1.0	0.0	153.0	109.000000	1880.000000	-37.83820	144.94690	1
								171.0	167.000000	1890.000000	-37.83850	144.94090	1
								184.0	100.000000	1905.000000	-37.83480	144.94430	1
								197.0	176.866248	1954.081176	-37.83430	144.93620	1
10	1	2	12.1	3083.0	10.0	3.0	2.0	313.0	176.866248	2006.000000	-37.71098	145.05381	2

Length: 13513, dtype: int64



After Oversampling

```
X, y = r.fit_resample(X,y)
```

```
X.shape, y.shape
```

```
((249052, 15), (249052,))
```

```
y.value_counts()
```

```
1480000.0    113
978500.0     113
920500.0     113
2633000.0    113
760500.0     113
...
2220000.0    113
667000.0     113
2105000.0    113
2177000.0    113
1323000.0    113
```

```
Name: Price, Length: 2204, dtype: int64
```

```
X.value_counts()
```


Rooms	Type	Method	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Latitude	Longitude	R
2	2	1	3.3	3141.0	2.0	2.0	2.0	17200.0	80.737121	2000.000000	-37.83610	144.99660	1
1	2	2	11.8	3204.0	1.0	1.0	1.0	0.0	80.737121	1980.016708	-37.90320	145.04550	1
2	2	2	9.1	3015.0	2.0	1.0	1.0	0.0	80.737121	1970.000000	-37.82880	144.87110	3
3	1	2	7.0	3013.0	3.0	1.0	1.0	248.0	176.866248	1954.081176	-37.81570	144.89250	3
4	1	2	7.7	3015.0	4.0	2.0	0.0	389.0	158.000000	1990.000000	-37.82840	144.88610	3
3	1	1	12.6	3020.0	3.0	1.0	1.0	700.0	145.000000	1960.000000	-37.79330	144.84110	3
4	1	1	15.5	3038.0	4.0	2.0	2.0	713.0	164.000000	1982.000000	-37.72305	144.81074	3
3	1	1	12.6	3020.0	3.0	1.0	1.0	500.0	176.866248	1954.081176	-37.78210	144.84560	3
								0.0	126.000000	1950.000000	-37.77960	144.84550	3
10	1	2	12.1	3083.0	10.0	3.0	2.0	313.0	176.866248	2006.000000	-37.71098	145.05381	2

Length: 13513, dtype: int64

▼ Step 2: Splitting the data

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=2529)#since test size is given as 30%, train size is 70
```

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((174336, 15), (74716, 15), (174336,), (74716,))
```

▼ Standard Scaling the dataset

```
from sklearn.preprocessing import StandardScaler
```

```
s=StandardScaler()
```

```
X_train_s=s.fit_transform(X_train)#Scaling train data
```

```
X_test_s=s.fit_transform(X_test)#Scaling test data
```

▼ Visualisation and Impact of Scaling

Reduced impact of outlier

```
X_train_s=pd.DataFrame(X_train_s,columns=X_train.columns)
```

```
X_test_s=pd.DataFrame(X_test_s,columns=X_test.columns)
```

```
fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))
ax1.scatter(X_train['YearBuilt'], X_train['Landsize'],color='green')
ax1.set_title("Before Scaling")
ax2.scatter (X_train_s['YearBuilt'], X_train_s['Landsize'],color="red")
ax2.set_title("After Scaling")
plt.show()
```

Trying this code on different attributes will help you understand the impact of scaling

▼ Step 3: Creating Model

```
#from sklearn.linear_model import LinearRegression
#model=LinearRegression()#Linear Regression Model
```

```
#from sklearn.neighbors import KNeighborsRegressor
```

```
#model=KNeighborsRegressor()#K-Nearest Neighbour Model
```

```
#from sklearn.tree import DecisionTreeRegressor  
#model=DecisionTreeRegressor()#Decision tree Model
```

```
from sklearn.ensemble import RandomForestRegressor  
model=RandomForestRegressor()#Random Forest Model
```

▼ Step 4: Training Model

```
model.fit(X_train_s,y_train)#model that is used in the previous step is being trained in this step
```

```
RandomForestRegressor()
```

▼ Step 5: Predicting Model

```
y_pred=model.predict(X_test_s)#Predicting the target for the given data
```

```
y_pred
```

```
array([1.93, 2. , 2. , ..., 2. , 1. , 1. ])
```

▼ Step 6: Accuracy

```
from sklearn.metrics import mean_absolute_percentage_error,mean_absolute_error,r2_score#checking for the accuracy of the prediction m
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.05907422347720855
```

```
mean_absolute_error(y_test,y_pred)
```

```
0.07926986688180719
```

```
r2_score(y_test,y_pred)
```


```
0.8529766602228644
```

▼ Sample Future Predictions Example

```
df_new=df.sample(1)#taking a sample set from the dataset
```

```
df_new#displaying sample set
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	Bu
8245	Prahran	303/10 Hillingdon PI	2	2	645000.0	1	Gary	2017- 05-13	4.5	3181.0	2.0	1.0	1.0	2842.0	



```
X_new=df_new.drop(['Suburb','Address','SellerG','Date','Price'],axis=1)#defining X or independent variables of samole set
```

```
X_new.shape
```

```
(1, 15)
```

```
y_pred_new=model.predict(X_new)#sample prediction
```

```
y_pred_new
```

```
array([2377000.])
```

Part 2: Classification Analysis

▼ Step 1: Define X and y

```
y=df['Type']# 'Type' is our Classification Target
```

```
y.shape
```

```
(13580,)
```

```
X=df.drop(['Suburb','Address','SellerG','Date','Type'],axis=1)
```

```
X.shape
```

```
(13580, 15)
```

▼ Dealing with Undersampling Data

```
|  
'Type'
```

```
from imblearn.under_sampling import RandomUnderSampler
```

```
r=RandomUnderSampler(random_state=2408)
```

```
X, y = r.fit_resample(X,y)
```

```
X.shape, y.shape
```

```
((8262, 15), (8262,))
```

```
y.value_counts()
```

```
1    4131
2    4131
Name: Type, dtype: int64
```

```
X.value_counts()
```

Rooms	Price	Method	Distance	Postcode	Bedroom2	Bathroom	Car	Landsize	BuildingArea	YearBuilt	Lattitude	Longtitu
2	965000.0	1	7.3	3146.0	2.0	2.0	1.0	704.0	140.046323	1998.988189	-37.85698	145.0468
	890500.0	2	2.4	3121.0	2.0	2.0	2.0	180.0	123.000000	1940.000000	-37.81443	144.9907
	435000.0	2	7.8	3124.0	2.0	1.0	1.0	896.0	77.000000	1960.000000	-37.84790	145.0958
4	1817000.0	2	4.2	3031.0	4.0	2.0	1.0	309.0	176.866248	1954.081176	-37.79100	144.9280
3	720000.0	2	7.8	3058.0	3.0	2.0	2.0	531.0	112.000000	2016.000000	-37.74240	144.9571
2	802000.0	1	5.4	3101.0	2.0	1.0	1.0	0.0	80.737121	1980.016708	-37.80456	145.0374
	801250.0	1	11.0	3018.0	2.0	1.0	1.0	245.0	80.737121	1980.016708	-37.87046	144.8348
	801000.0	1	14.5	3188.0	2.0	1.0	1.0	217.0	107.000000	1994.000000	-37.94320	145.0348
			6.9	3039.0	2.0	1.0	0.0	133.0	176.866248	1954.081176	-37.77120	144.9254
8	2950000.0	1	11.0	3147.0	9.0	7.0	4.0	1472.0	618.000000	2009.000000	-37.87290	145.0788

Length: 8251, dtype: int64

▼ Step 2: Splitting the data

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,stratify=y,random_state=2408)
```

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((5783, 15), (2479, 15), (5783,), (2479,))
```

▼ Standard Scaling the dataset

```
df.describe()
```

	Rooms	Type	Price	Method	Distance	Postcode	Bedroom2	Bathroom	Car
count	13580.000000	13580.000000	1.358000e+04	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000	13580.000000

```

from sklearn.preprocessing import StandardScaler

s=StandardScaler()
.....
X_train_s=s.fit_transform(X_train)
50% 3.000000 1.000000 9.030000e+05 1.000000 9.200000 3084.000000 3.000000 1.000000 2.000000
X_test_s=s.fit_transform(X_test)

```

▼ Visualisation and Impact of Scaling

Reduced impact of outlier

```

X_train_s=pd.DataFrame(X_train_s,columns=X_train.columns)
X_test_s=pd.DataFrame(X_test_s,columns=X_test.columns)

fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(12, 5))

ax1.scatter(X_train['YearBuilt'], X_train['Price'],color='black')
ax1.set_title("Before Scaling")
ax2.scatter (X_train_s['YearBuilt'], X_train_s['Price'],color="red")
ax2.set_title("After Scaling")
plt.show()

```

▼ Step 3: Creating Model


```
#from sklearn.linear_model import LogisticRegression
#model=LogisticRegression()#Logistic Regression Model

#from sklearn.neighbors import KNeighborsClassifier
#model=KNeighborsClassifier()#K-Nearest Neighbour Model

#from sklearn.tree import DecisionTreeClassifier
#model=DecisionTreeClassifier()#Decision Tree Model

#from sklearn.svm import SVC
#model=SVC() #Support Vector Machine

from sklearn.ensemble import RandomForestClassifier
model=RandomForestClassifier()
```

▼ Step 4: Training Model

```
model.fit(X_train_s,y_train)

RandomForestClassifier()
```

▼ Step 5: Predicting Model

```
y_pred=model.predict(X_test_s)

y_pred
```

```
array([2, 2, 2, ..., 2, 1, 1])
```

▼ Step 6: Accuracy

```
from sklearn.metrics import confusion_matrix, classification_report
```

confusion_matrix(y_test,y_pred)#confusion matrix is the best way to represent the accuracy of a Classification Model. It represents T

```
array([[1178, 61],
       [ 37, 1203]])
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
1	0.97	0.95	0.96	1239
2	0.95	0.97	0.96	1240
accuracy			0.96	2479
macro avg	0.96	0.96	0.96	2479
weighted avg	0.96	0.96	0.96	2479

▼ Future Predictions Example

```
df_new=df.sample(1)#sample set
```

```
df_new
```

	Suburb	Address	Rooms	Type	Price	Method	SellerG	Date	Distance	Postcode	Bedroom2	Bathroom	Car	Landsi
11596	Bentleigh	12 Wood St	5	1	1500000.0	2	hockingstuart	2017-07-22	11.4	3204.0	5.0	2.0	2.0	591



```
X_new=df_new.drop(['Suburb','Address','Type','SellerG','Date'],axis=1)
```

```
X_new.shape
```

```
(1, 15)
```

```
y_pred_new=model.predict(X_new)#sample prediction.
```

```
y_pred_new
```

```
array([1])
```

Link of the Colab file:

<https://colab.research.google.com/drive/1nZuh04PYrE9kxCzCDbyYsH-b3czkKgpD?usp=sharing>

