

Mini Project 7: Exchange Vehicle Price Prediction

1. Bike Price Prediction

▼ Step 1: Import library

```
import pandas as pd
```

```
import numpy as np
```

```
import seaborn as sns
```

▼ Step 2: Import Data

```
df=pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/Bike%20Prices.csv')
```

```
df
```

Saved successfully!



			Model	Selling_Price	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
0	TVS	TVS XL 100		30000	2017	Individual	1st owner	8000	30490.0
1	Bajaj	Bajaj ct 100		18000	2017	Individual	1st owner	35000	32000.0
2	Yo	Yo Style		20000	2011	Individual	1st owner	10000	37675.0
3	Bajaj	Bajaj Discover 100		25000	2010	Individual	1st owner	43000	42859.0
4	Bajaj	Bajaj Discover 100		24999	2012	Individual	2nd owner	35000	42859.0
...
1056	Royal	Royal Enfield Electra 5 S		90000	2012	Individual	1st owner	40000	NaN
1057	Hero	Hero Honda Hunk		20000	2010	Individual	1st owner	17000	NaN
1058	Bajaj	Bajaj Pulsar 220 DTS-i		60000	2014	Individual	1st owner	16000	NaN

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1061 entries, 0 to 1060
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Brand                 1061 non-null  object  
1   Model                 1061 non-null  object  
2   Selling_Price         1061 non-null  int64   
3   Year                  1061 non-null  int64   
4   Seller_Type           1061 non-null  object  
5   Owner                 1061 non-null  object  
6   KM_Driven             1061 non-null  int64   
7   Ex_Showroom_Price     626 non-null   float64  
dtypes: float64(1), int64(3), object(4)
memory usage: 66.4+ KB

```

▼ Auxillary Step: Drop Missing Value Rows

Saved successfully!



```
df = df.dropna()
```

```
df.describe()
```

	Selling_Price	Year	KM_Driven	Ex_Showroom_Price
count	626.000000	626.000000	626.000000	6.260000e+02
mean	59445.164537	2014.800319	32671.576677	8.795871e+04
std	59904.350888	3.018885	45479.661039	7.749659e+04
min	6000.000000	2001.000000	380.000000	3.049000e+04
25%	30000.000000	2013.000000	13031.250000	5.485200e+04
50%	45000.000000	2015.000000	25000.000000	7.275250e+04
75%	65000.000000	2017.000000	40000.000000	8.703150e+04
max	760000.000000	2020.000000	585659.000000	1.278000e+06

```
df.shape
```

```
(626, 8)
```

```
df.columns
```

```
Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner',  
      'KM_Driven', 'Ex_Showroom_Price'],  
      dtype='object')
```

▼ Auxillary Step: Fetching Categories and Counts of y

Saved successfully!



```

Brand
Honda      170
Bajaj      143
Hero       108
Yamaha     94
Royal      40
TVS        23
Suzuki     18
KTM        6
Mahindra   6
Kawasaki   4
UM         3
Activa     3
Harley     2
Vespa      2
BMW        1
Hyosung    1
Benelli    1
Yo         1
dtype: int64

```

```
df[['Model']].value_counts()
```

```

Model
Honda Activa [2000-2015]      23
Honda CB Hornet 160R         22
Bajaj Pulsar 180             20
Yamaha FZ S V 2.0            16
Bajaj Discover 125           16
..
Royal Enfield Thunderbird 500 1
Royal Enfield Continental GT [2013 - 2018] 1
Royal Enfield Classic Stealth Black 1
Royal Enfield Classic Squadron Blue 1
Yo Style                     1
Length: 183, dtype: int64

```

Saved successfully!

```
Seller_Type
Individual    623
Dealer        3
dtype: int64
```

```
df[['Owner']].value_counts()
```

```
Owner
1st owner    556
2nd owner    66
3rd owner     3
4th owner     1
dtype: int64
```

▼ Auxillary Step: Encoding Categorical Features

```
df.replace({'Seller_Type':{'Individual':0,'Dealer':1}},inplace=True)
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self._setitem_single_column(ilocs[0], value, pi)
```

```
df.replace({'Owner':{'1st owner':0,'2nd owner':1,'3rd owner':2,'4th owner':3}},inplace=True)
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py:1773: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

Saved successfully!

ation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
locs[0], value, pi)

▼ Step 3. Define y and X

```
y = df['Selling_Price']
```

```
y.shape
```

```
(626,)
```

```
y
```

```
0      30000  
1      18000  
2      20000  
3      25000  
4      24999
```

```
...
```

```
621    330000  
622    300000  
623    425000  
624    760000  
625    750000
```

```
Name: Selling_Price, Length: 626, dtype: int64
```

```
#X=df.drop(['Brand', 'Model', 'Selling_Price'],axis=1)
```

```
#or
```

```
X=df[['Year', 'Seller_Type', 'Owner', 'KM_Driven', 'Ex_Showroom_Price']]
```

```
X.shape
```

Saved successfully!



X

	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
0	2017	0	0	8000	30490.0
1	2017	0	0	35000	32000.0
2	2011	0	0	10000	37675.0
3	2010	0	0	43000	42859.0
4	2012	0	1	35000	42859.0
...
621	2014	0	3	6500	534000.0
622	2011	0	0	12000	589000.0
623	2017	0	1	13600	599000.0
624	2019	0	0	2800	752020.0
625	2013	0	1	12000	1278000.0

626 rows × 5 columns

▼ Step 4: Splitting Data

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_state=2529)
```

Saved successfully!



hape,y_test.shape

```
((438, 5), (188, 5), (438,), (188,))
```

▼ Step 5: Creating Model

```
from sklearn.linear_model import LinearRegression
model=LinearRegression()
```

▼ Step 6: Training Model

```
model.fit(X_train,y_train)
```

```
LinearRegression()
```

▼ Step 7: Predicting Model

```
y_pred=model.predict(X_test)
```

```
y_pred.shape
```

```
(188,)
```

```
y_pred
```

```
array([ 27210.52271465,  56340.08335163,  63471.94671996,  53627.63844785,
        55612.75744268,  53888.92259719,  33751.35275102,  60311.4950183 ,
```


Saved successfully!



```
49332954, 27826.7399381 , 49919.83255841,
12664064, 48277.75426038, 127646.56079335,
70047.10661635, 39350.67963653, 36081.03597878, 45360.79436339,
48079.89470577, 44803.02464799, 55161.44026111, 71041.51821318,
91689.22699159, 49301.53594645, 55988.19326252, 108171.54600296,
32771.06897901, 25468.20072996, 17128.61806164, 179271.41130746,
45698.99857622, 31371.09285079, 67886.52106737, 41492.49575815,
56855.22238602, 47820.47003468, 74682.14053958, 24984.21822736,
55374.00513699, 41412.36775222, 67991.60287764, 26553.59421844,
89788.69870689, 45764.83633686, 133888.03770389, 106988.113825 ,
71176.40667714, 25332.25485946, 79512.43778826, 63914.38088173,
28632.12110986, 53656.13623937, -5396.37132904, 70377.44571174,
33313.03576476, 53994.92478411, 67509.85836352, 59735.05378847,
22199.83644217, 15374.18984158, 44510.76819427, 30279.52476752,
108243.77037514, 19291.8895874 , 53614.312976 , 59230.23269131,
60174.2108109 , 45924.63468736, 25770.81883496, 63471.36257814,
242123.45729792, 61387.72544548, 56510.98127074, 48123.28087213,
51668.27442011, 90279.76190495, 14827.76533556, 112437.70820504,
35066.88027405, 30902.41069172, 31441.48921433, 125593.75847157,
27705.38813164, -11590.29205553, 15582.17108685, 75113.64511232,
504085.44522282, 123545.42050116, 74770.89327697, 50747.47663245,
44174.3618212 , 25426.7156106 , 30298.3052462 , 47625.67836414,
27850.37544807, 28845.23330928, 31580.38624692, 32309.63375635,
47979.16788554, 65955.46375944, 13432.28218017, 15368.80064986,
31973.23052409, 110353.92870546, 68181.49509136, 23143.49139797,
53194.65732076, 34603.36376989, 56002.50967868, 62432.66994305,
391470.77533201, 3558.29480891, 36019.18494305, 70876.34866549,
72890.00667025, 137596.01384364, 27620.36308877, 135789.30486854,
39674.40366791, 58367.0924453 , 42401.21202624, 61864.4379567 ,
42688.89652842, 63710.34571021, 10604.39360071, 38458.82820943,
112251.84744225, 115403.00577536, 13658.41734785, 36196.83359584,
54146.22998932, 97297.85724851, 55029.68137265, 22923.26533437,
104569.97029689, 41965.75852017, 38759.68546491, 28930.61369011,
45231.66612551, 48475.43422775, 26739.7225731 , 53598.65972203,
32558.54954524, 32212.22834942, 68172.98738422, 71839.47716461,
32003.46692215, 40652.69995971, 39935.92211843, 63444.41846202,
44545.5818771 , 120873.38389616, 60926.58683174, 62641.82167496,
60816.47379994, 27098.95433573, 26803.64749618, 48956.00468627,
62032.88118713, 26471.97495723, 104937.23068766, 132903.3578847 ,
37469.2040942 , 57579.12080094, 40371.00915736, -7039.40662503,
42554145, 52153.21149321, 56453.74542453,
```

Saved successfully!



```
46870273, 49505.97985573, 24288.36959514,  
26333955, 23399.66596746, 63678.40865459,  
70144.29372668, 33434.89010059, 60885.29444481, 58389.55370878,  
35118.7040348 , 58729.4540196 , 34627.9532246 , 38583.4623973 ])
```

▼ Step 8: Accuracy

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
mean_squared_error(y_test, y_pred)
```

```
554715615.5043668
```

```
mean_absolute_error(y_test, y_pred)
```

```
12225.7370104107
```

```
r2_score(y_test, y_pred)
```

```
0.8810414402984937
```


▼ Data Visualisation

```
import matplotlib.pyplot as plt  
plt.scatter(y_test, y_pred)  
plt.xlabel('Actual Prices')  
plt.ylabel('Predicted Prices')  
plt.title('Actual v/s Predicted Prices')  
plt.show
```

Saved successfully!



W>



▼ Step 10: Future Predictions Example

```
df_new=df.sample(1)
```

```
df_new
```

	Brand	Model	Selling_Price	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Price
367	Yamaha	Yamaha FZ16	30000	2009	0	1	26000	78712.0

```
X_new=df_new.drop(['Brand', 'Model', 'Selling_Price'],axis=1)
```

```
X_new.shape
```

```
(1, 5)
```

Saved successfully!

`y_pred_new``array([21810.4089662])`

2. Car Price Prediction

▼ Step 1: Import library

`import pandas as pd``import numpy as np``import seaborn as sns`

▼ Step 2: Import Data

`df=pd.read_csv('https://github.com/YBI-Foundation/Dataset/raw/main/Car%20Price.csv')``df`

Saved successfully!



			Model	Year	Selling_Price	KM_Driven	Fuel	Seller_Type	Transmission	Owner
0	Maruti		Maruti 800 AC	2007	60000	70000	Petrol	Individual	Manual	First Owner
1	Maruti		Maruti Wagon R LXI Minor	2007	135000	50000	Petrol	Individual	Manual	First Owner
2	Hyundai		Hyundai Verna 1.6 SX	2012	600000	100000	Diesel	Individual	Manual	First Owner
3	Datsun		Datsun RediGO T Option	2017	250000	46000	Petrol	Individual	Manual	First Owner
4	Honda		Honda Amaze VX i-DTEC	2014	450000	141000	Diesel	Individual	Manual	Second Owner
...
4335	Hyundai		Hyundai i20 Magna 1.4 CRDi (Diesel)	2014	409999	80000	Diesel	Individual	Manual	Second Owner
4336	Hyundai		Hyundai i20 Magna 1.4 CRDi	2014	409999	80000	Diesel	Individual	Manual	Second Owner
4337	Maruti		Maruti 800 AC BSIII	2009	110000	83000	Petrol	Individual	Manual	Second Owner
4338	Hyundai		Hyundai Creta 1.6 CRDi SX Option	2016	865000	90000	Diesel	Individual	Manual	First Owner

df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4340 entries, 0 to 4339
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Brand            4340 non-null   object
1   Model            4340 non-null   object
2   Year             4340 non-null   int64
3   Selling_Price    4340 non-null   int64
4   KM_Driven        4340 non-null   int64
5   Fuel             4340 non-null   object
6   Seller_Type      4340 non-null   object
7   Transmission     4340 non-null   object
8   Owner            4340 non-null   object
dtypes: int64(3), object(6)
memory usage: 305.3+ KB

```

Saved successfully!



	Year	Selling_Price	KM_Driven
count	4340.000000	4.340000e+03	4340.000000
mean	2013.090783	5.041273e+05	66215.777419
std	4.215344	5.785487e+05	46644.102194
min	1992.000000	2.000000e+04	1.000000
25%	2011.000000	2.087498e+05	35000.000000
50%	2014.000000	3.500000e+05	60000.000000
75%	2016.000000	6.000000e+05	90000.000000
max	2020.000000	8.900000e+06	806599.000000

```
df.shape
```

```
(4340, 9)
```

```
df.columns
```

```
Index(['Brand', 'Model', 'Year', 'Selling_Price', 'KM_Driven', 'Fuel',
      'Seller_Type', 'Transmission', 'Owner'],
      dtype='object')
```

▼ Auxillary Step: Fetching Categories and Counts of y

```
df[['Brand']].value_counts()
```

```
Brand
```

```
Model 4340
```

Saved successfully!



Tata	361
Honda	252
Ford	238
Toyota	206
Chevrolet	188
Renault	146
Volkswagen	107
Skoda	68
Nissan	64
Audi	60
BMW	39
Fiat	37
Datsun	37
Mercedes-Benz	35
Mitsubishi	6
Jaguar	6
Land	5
Ambassador	4
Volvo	4
Jeep	3
OpelCorsa	2
MG	2
Isuzu	1
Force	1
Daewoo	1
Kia	1

dtype: int64

```
df[['Model']].value_counts()
```

Model	
Maruti Swift Dzire VDI	69
Maruti Alto 800 LXI	59
Maruti Alto LXI	47
Hyundai EON Era Plus	35
Maruti Alto LX	35
..	
1	

Saved successfully!



AW

1
1
1
1

Mahindra KUV 100 mFALCON G80 K2 Plus

Volvo XC60 D5 Inscription

Length: 1491, dtype: int64

```
df[['Fuel']].value_counts()
```

Fuel

Diesel 2153

Petrol 2123

CNG 40

LPG 23

Electric 1

dtype: int64

```
df[['Seller_Type']].value_counts()
```

Seller_Type

Individual 3244

Dealer 994

Trustmark Dealer 102

dtype: int64

```
df[['Transmission']].value_counts()
```

Transmission

Manual 3892

Automatic 448

dtype: int64

```
df[['Owner']].value_counts()
```

Owner

First Owner 2832

Second Owner 1106

Saved successfully!



dtype: int64

▼ Auxillary Step: Encoding Categorical Features

```
df.replace({'Fuel':{'Petrol':0,'Diesel':1,'CNG':2,'LPG':3,'Electric':4}},inplace=True)
```

```
df.replace({'Seller_Type':{'Individual':0,'Dealer':1,'Trustmark Dealer':2}},inplace=True)
```

```
df.replace({'Transmission':{'Manual':0,'Automatic':1}},inplace=True)
```

```
df.replace({'Owner':{'First Owner':0,'Second Owner':1,'Third Owner':2,'Fourth & Above Owner':3,'Test Drive Car':4}},inplace=True)
```

▼ Step 3. Define y and X

```
y = df['Selling_Price']
```

```
y.shape
```

```
(4340,)
```

```
y
```

```
0      60000
1     135000
2     600000
```

Saved successfully!



```
...
4335    409999
4336    409999
4337    110000
4338    865000
4339    225000
Name: Selling_Price, Length: 4340, dtype: int64
```

```
#X=df.drop(['Brand', 'Model', 'Selling_Price'],axis=1)
#or
X=df[['Year', 'KM_Driven','Fuel', 'Seller_Type','Transmission', 'Owner']]
```

```
X.shape
```

```
(4340, 6)
```

```
X
```

Saved successfully!



	Year	Price	Engine	Transmission	Owner
0	2007	70000	0	0	0
1	2007	50000	0	0	0
2	2010	100000	1	0	0

▼ Step 4: Splitting Data

```
from sklearn.model_selection import train_test_split
```

	Year	Price	Engine	Transmission	Owner
4335	2014	80000	1	0	1

```
X_train,X_test,y_train,y_test=train_test_split(X,y,train_size=0.7,random_state=2529)
```

	Year	Price	Engine	Transmission	Owner
4337	2000	80000	0	0	1

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((3038, 6), (1302, 6), (3038,), (1302,))
```

▼ Step 5: Creating Model

```
from sklearn.linear_model import LinearRegression
model=LinearRegression()
```

▼ Step 6: Training Model

```
model.fit(X_train,y_train)
```

```
LinearRegression()
```

Saved successfully!



```
y_pred=model.predict(X_test)
```

```
y_pred.shape
```

```
(1302,)
```

```
y_pred
```

```
array([502458.82786413, 646333.17428704, 521962.74075836, ...,  
       620183.32683781, 315403.8278857 , 731862.54196037])
```

▼ Step 8: Accuracy

```
from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
```

```
mean_squared_error(y_test,y_pred)
```

```
193242972302.19553
```

```
mean_absolute_error(y_test,y_pred)
```

```
228808.95522977872
```

```
r2_score(y_test,y_pred)
```


Saved successfully!



▼ Data Visualisation

```
import matplotlib.pyplot as plt
plt.scatter(y_test,y_pred)
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual v/s Predicted Prices')
plt.show
```

<function matplotlib.pyplot.show>



▼ Future Predictions Example

```
df_new=df.sample(1)
```

Saved successfully!



	Brand	Model	Year	Selling_Price	KM_Driven	Fuel	Seller_Type	Transmission	Owner
306	Mahindra	Mahindra Bolero Power Plus AC BSIV PS	2015	430000	200000	1	0	0	0

```
X_new=df_new.drop(['Brand', 'Model', 'Selling_Price'],axis=1)
```

```
X_new.shape
```

```
(1, 6)
```

```
y_pred_new=model.predict(X_new)
```

```
y_pred_new
```

```
array([481052.59195395])
```

Link for the same:

<https://colab.research.google.com/drive/1lct9DHDz1BwNsTbQrk1aHJU3ki1x1cIN?usp=sharing>

Saved successfully!

